

Cupcast III: The 2026 World Cup Forecast

Carlos Garcia

June 2026

Abstract

The forecast: advancement probabilities for all 48 teams and the podium distribution from 50,000 Monte Carlo simulations of the full 2026 FIFA World Cup under the Dixon–Coles model with squad-quality shrinkage [1] specified in Paper I and validated in Paper II, together with path-dependence and key-player sensitivity analyses and a comparison against bookmaker-implied probabilities. Per-match predictions for all 104 matches accompany this paper as machine-readable artifacts in the repository’s `outputs/` directory.

1 Headline Forecast

Table positions are probabilities of reaching each stage, from 50,000 seeded simulations (every number in this paper regenerates from `python -m cupcast.report.build`):

Team	R32	QF	SF	Final	Champion
Argentina	0.980	0.519	0.370	0.248	0.163
Spain	0.992	0.477	0.351	0.240	0.150
Brazil	0.983	0.456	0.292	0.163	0.093
England	0.979	0.446	0.269	0.154	0.084
France	0.942	0.412	0.250	0.134	0.071
Portugal	0.936	0.402	0.229	0.128	0.065
Colombia	0.915	0.349	0.189	0.101	0.050
Germany	0.975	0.350	0.191	0.090	0.040
Belgium	0.951	0.372	0.166	0.085	0.039
Netherlands	0.899	0.305	0.163	0.077	0.035
Morocco	0.915	0.305	0.159	0.074	0.033
Ecuador	0.935	0.262	0.130	0.056	0.024

2 The Podium

Ranking marginal probabilities per podium step gives the headline pick; the joint distribution below accounts for the fact that one team’s podium finish constrains the others’ (the third step is decided by the actual third-place play-off in every simulation):

Champion	Runner-up	Third	Probability
Argentina	Spain	Brazil	0.0045
Spain	Argentina	Brazil	0.0038
Argentina	Spain	England	0.0032
Argentina	Spain	France	0.0031
Spain	Argentina	France	0.0026
Spain	Argentina	England	0.0026
Spain	England	Argentina	0.0026
Spain	Brazil	Argentina	0.0025

3 Path Dependence

Group-stage outcomes reshape each contender’s bracket path. The conditional championship probabilities below are computed within the same simulation set; the gap between finishing first and third in the group is the price of the harder Round-of-32 draw under the official bracket:

Team	P(champ)	given 1st	given 2nd	given 3rd q.
Spain	0.1500	0.1571	0.1329	0.1298
Argentina	0.1629	0.1720	0.1452	0.1684
England	0.0840	0.0923	0.0710	0.0827
Japan	0.0183	0.0252	0.0186	0.0223
France	0.0709	0.0826	0.0621	0.0730
Brazil	0.0926	0.0925	0.0993	0.0906
United States	0.0037	0.0061	0.0046	0.0047
Mexico	0.0082	0.0094	0.0085	0.0082

4 Key-Player Sensitivity

Each scenario removes or limits a named player, reallocates his expected minutes to the next players in the position group, rebuilds the squad composite and shrinkage prior of Paper I §6, and reruns the identical seeded simulation:

Scenario	Before	After	Δ champion	Δ SF
Rodri out	0.1500	0.1495	-0.0005	0.0031
Lamine Yamal out	0.1500	0.1473	-0.0027	-0.0023
Bellingham at 70 percent	0.0840	0.0829	-0.0010	-0.0044
Messi limited to 60 minutes	0.1629	0.1627	-0.0002	0.0002
Mbappé out	0.0709	0.0707	-0.0002	0.0008
Vinícius Júnior out	0.0926	0.0935	0.0009	0.0014

No scenario moves a championship probability beyond the Monte Carlo noise floor ($\approx \pm 0.3\text{pp}$ at 50,000 simulations). This is a structural property worth stating plainly: with ratings dominated by match results and the tuned light shrinkage of Paper I, a single player’s availability shifts the squad composite, the composite shifts the prior, and the prior carries only $\sim 20\%$ weight for a data-rich team — each link damps the effect. The model’s genuine sensitivity is path-dependence (the previous section), not single-player availability; a reader who believes individual stars matter more than this should treat the headline probabilities as conditional on roughly current squad quality.

5 Model versus Market

Consensus bookmaker championship probabilities (outright odds across books, proportional overround removal [2]) against the model, sorted by absolute disagreement:

Team	Model	Market	Edge
Argentina	0.163	0.080	0.082
France	0.071	0.146	-0.075
Colombia	0.050	0.020	0.031
Portugal	0.065	0.096	-0.031
England	0.084	0.105	-0.021
Morocco	0.033	0.016	0.018
Belgium	0.039	0.022	0.016
Germany	0.040	0.056	-0.015
Ecuador	0.024	0.009	0.015
Norway	0.016	0.027	-0.011
Brazil	0.093	0.082	0.010
Uruguay	0.020	0.011	0.009
Netherlands	0.035	0.043	-0.008
Croatia	0.015	0.008	0.007
Switzerland	0.019	0.012	0.006

Where the model and the market disagree materially, the burden of proof sits with the model; the largest edges are discussed alongside the executive summary in `outputs/`.

6 Limitations

Expected minutes use a caps-based depth heuristic (Paper I §5) — researched likely line-ups for the main contenders are the highest-value refinement. Fair-play tiebreakers are replaced by lots; rest days and travel distance are not modelled; squad composites cover only players in the big five European leagues, with Elo carrying the remainder. The market comparison reflects a single pre-tournament odds snapshot.

References

- [1] M. J. Dixon and S. G. Coles. Modelling association football scores and inefficiencies in the football betting market. *Journal of the Royal Statistical Society: Series C (Applied Statistics)*, 46(2):265–280, 1997. doi: 10.1111/1467-9876.00065.
- [2] The Odds API. Sports odds api. <https://the-odds-api.com>. Accessed June 2026.